


Highly-scalable Concurrent Objects

Deepthi Akkoorath¹ Annette Bieniusa¹

¹University of Kaiserslautern
AG Software Technology

18. Apr 2016


Concurrent Shared-Memory Systems



Correctness = Linearizability \implies Synchronization
Scalable?

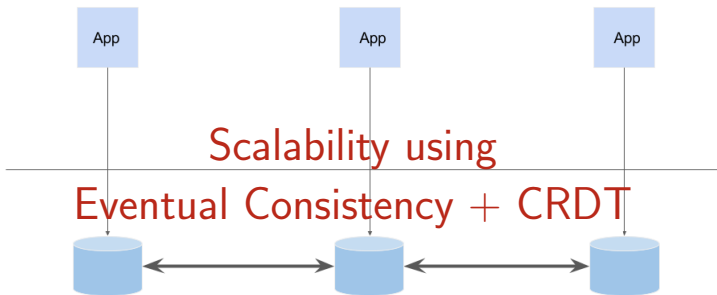
Shared Objects

Geo-Replicated Systems



Shared Objects

Geo-Replicated Systems



Shared Objects

Shared Memory Systems / Multi-core

- Many Scalable Data-structures exists
- Lock-free data-structures using Transactional Memory
- Scalable under contention?

Shared Objects

Shared Memory Systems / Multi-core


- Many Scalable Data-structures exists
- Lock-free data-structures using Transactional Memory
- Scalable under contention?

Weakening Linearizability \implies Scalability?


Mergeable Objects in Shared Memory Systems

Thread 1


Thread 2




Mergeable Objects in Shared Memory Systems



Mergeable Objects in Shared Memory Systems




Mergeable Objects in Shared Memory Systems



Mergeable Transactions

Mergeable Objects in STM



CRDTs as Mergeable Objects

- Example: G-Set
 - a Set
 - Merge = union of two sets
- Example: G-Counter
 - A vector
 - Merge = max of each element in the vector

CRDTs as Mergeable Objects


- Example: G-Set
 - a Set
 - Merge = union of two sets
- Example: G-Counter
 - A vector
 - Merge = max of each element in the vector

CRDT Merge is Expensive

Properties of Mergeable Objects

Persistence


- Multiple Versions co-exists
 - Local Versions
 - Global Version
- A **data structure is persistent** if multiple versions of it are accessible




Properties of Mergeable Objects

Persistence

- Multiple Versions co-exists
 - Local Versions
 - Global Version
- A **data structure** is **persistent** if multiple versions of it are accessible
- Copy entire Object?



Persistence in List




Mergeability

- Semantic Mergeability
 - Exploit object semantics (e.g. idempotence, commutativity) to obtain merge operation
 - Similar to CRDTs [Shapiro et. al '12]
- Structural Mergeability
 - How efficiently two versions can be merged
 - Share (parts of) the data structure between concurrent threads
 - Merge by modifying the data structure


Mergeable Counter

- 2 Integers: (G, L)
- increment = $(G, L+1)$
- merge $(G, 0) (G', L) = (G+L, 0)$
- Persistence by “Copy object”
- Semantic Merge by arithmetic properties


Add only Bag



After $T1$ commits:




After $T2$ commits:



OR-Set


Binary Search Tree



- Node = (Val, Version info)
- Version info = $\langle first, last, removed \rangle$
-

OR-Set


Binary Search Tree



- Node = (Val, Version info)
- Version info = $\langle first, last, removed \rangle$
- Lookup

OR-Set


Binary Search Tree



- Node = (Val, Version info)
- Version info = $\langle first, last, removed \rangle$
- Add

OR-Set


Binary Search Tree



- Node = (Val, Version info)
- Version info = $\langle first, last, removed \rangle$
- Merging Adds

OR-Set


Binary Search Tree



- Node = (Val, Version info)
- Version info = $\langle first, last, removed \rangle$
- Remove

OR-Set

Binary Search Tree



- Node = (Val, Version info)
- Version info = $\langle first, last, removed \rangle$
- Merging Removes

Conclusion

- Mergeable Objects analogous to CRDTs
 - Non-linearizable
 - Meaningful merge semantics
- Properties for efficiency
 - Persistence
 - Mergeability
- Scalable?



akkoorath@cs.uni-kl.de